

# Image descriptors


Diane Lingrand

Polytech SI4

2018 - 2019



?



- Bag of Words / Bag of Features
  - using SIFT, SURF, ...
- HOG
- Deep features

# Very basic idea : image as a pixel array




0.alexandre.hiltcher.009.txt

```
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 11,
30, 30, 11, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 30, 30,
30, 227, 30, 30, 0, 0, 0, 0, 0, 0, 0, 0, 3, 30, 30,
30, 30, 255, 227, 30, 30, 223, 0, 0, 0, 0, 0, 19,
171, 83, 30, 255, 255, 255, 255, 30, 255, 0, 0, 0,
0, 0, 0, 0, 30, 30, 255, 255, 255, 141, 0, 0, 255,
30, 0, 0, 0, 0, 0, 4, 227, 227, 237, 255, 30, 0, 0,
32, 255, 227, 0, 0, 0, 0, 0, 30, 30, 227, 30, 255,
255, 0, 0, 33, 255, 255, 0, 0, 0, 0, 0, 30, 104,
255, 171, 114, 255, 171, 255, 255, 255, 30, 0, 0,
0, 0, 1, 2, 30, 30, 118, 171, 255, 171, 202, 255, 58, 255, 0, 0, 0, 0, 0, 0,
27, 30, 30, 255, 255, 171, 30, 255, 58, 0, 0, 0, 0, 0, 0, 26, 30, 30, 30,
83, 41, 244, 55, 255, 0, 0, 0, 0, 0, 0, 0, 0, 149, 114, 255, 30, 30, 171, 104,
0, 0, 0, 0, 0, 0, 0, 0, 0, 30, 223, 255, 30, 30, 255, 0, 0, 0, 0, 0, 0, 0, 0,
0, 30, 181, 255, 255, 84, 30, 0, 0, 0, 0, 0]
```

## SIFT = Scale Invariant Feature Transform

- Detector
  - multi-scale
  - DOG laplacian (Difference Of Gaussians)
- Descriptor
  - edges orientations in the neighborhood

# SIFT descriptor




- vectors of 128 integers
- 4 steps :
  - interest points detection
  - gradients orientation in the neighborhood (16x16 pixels divided in 16 blocks of size 4x4)
  - orientation histogram (quantified on 8 values in blocks of 4x4 pixels)
    - $8 \times 4 \times 4 = 128$
  - normalisation

# SIFT descriptor



0.alexandre.hiltcher.006.png  
252 descriptors




0.alexandre.hiltcher.009.png  
182 descriptors


# SURF descriptor (Bay et al 2006)

SURF = Speeded-Up Robust Features

default size : 64





0.alexandre.hiltcher.006.png  
184 descriptors



0.alexandre.hiltcher.009.png  
121 descriptors




# Bag Of Words (BOW)



## HOG : Histogram of Gradients

- gradient computation (  $[-1 \ 0 \ 1]$  et  $[-1 \ 0 \ 1]^T$  )
- histogram construction
  - squared cells (from  $4 \times 4$  to  $12 \times 12$  pixels)
  - discretisation on 9 angle values
  - pixel votes proportional to gradient amplitude
- blocks construction
  - 1 block = several cells
  - normalisation of blocks
- HOG = concatenation of histograms


*Next slides from presentation by Seeman.*



0.alexandre.hiltcher.009.png



# Descriptor

1. Compute gradients on an image region of  $64 \times 128$  pixels
2. Compute histograms on 'cells' of typically  $8 \times 8$  pixels (i.e.  $8 \times 16$  cells)
3. Normalize histograms within overlapping blocks of cells (typically  $2 \times 2$  cells, i.e.  $7 \times 15$  blocks)
4. Concatenate histograms




# Cell histograms

- 9 bins for gradient orientations (0-180 degrees)
- Filled with magnitudes
- Interpolated trilinearly:
  - Bilinearly into spatial cells
  - Linearly into orientation bins






# Final Descriptor

- Concatenation of Blocks



- Visualization:



# Deep descriptor from GoogLeNet Inception v1

- already trained CNN :
  - try *GoogLe Net* by downloading the CNN topology<sup>1</sup>, the learned weights<sup>2</sup> and categories<sup>3</sup>
  - and the OpenCV example<sup>4</sup>
- output of convolutional layer [inception\\_5b/pool\\_proj](#)
  - Using code :

```
net.forward(outputBlobs,"inception_5b/pool_proj");
```
- 1 image  $\leftrightarrow$  1 vector of 6272 floats

---

1. [http://dl.caffe.berkeleyvision.org/bvlc\\_googlenet.caffemodel](http://dl.caffe.berkeleyvision.org/bvlc_googlenet.caffemodel)

2. [https://raw.githubusercontent.com/ludvix/opencv\\_contrib/master/modules/dnn/samples/bvlc\\_googlenet.prototxt](https://raw.githubusercontent.com/ludvix/opencv_contrib/master/modules/dnn/samples/bvlc_googlenet.prototxt)

3. [https://raw.githubusercontent.com/ludvix/opencv\\_contrib/master/modules/dnn/samples/synset\\_words.txt](https://raw.githubusercontent.com/ludvix/opencv_contrib/master/modules/dnn/samples/synset_words.txt)

4. [https://docs.opencv.org/3.2.0/d5/de7/tutorial\\_dnn\\_googlenet.html](https://docs.opencv.org/3.2.0/d5/de7/tutorial_dnn_googlenet.html)

- from an already trained CNN : Inception v3, VGG Net, ...
  - output of convolution layer (the last or the last before)
  - reduced size after convolution layers
- autoencoder
  - data specific encoding